

# Tensorflow Project

Let's take a look at the effectiveness of Neural Nets.

We'll use the Bank Authentication Data Set from the UCI repository.

The data consists of 5 columns:

- variance of Wavelet Transformed image (continuous)
- skewness of Wavelet Transformed image (continuous)
- curtosis of Wavelet Transformed image (continuous)
- entropy of image (continuous)
- class (integer)

Where class indicates whether or not a Bank Note was authentic.

```
In [1]: 1 import pandas as pd
```

```
In [21]: 1 data = pd.read_csv('bank_note_data.csv')
        2 data.head()
```

```
Out[21]:
```

	Image.Var	Image.Skew	Image.Curt	Entropy	Class
0	3.62160	8.6661	-2.8073	-0.44699	0
1	4.54590	8.1674	-2.4586	-1.46210	0
2	3.86600	-2.6383	1.9242	0.10645	0
3	3.45660	9.5228	-4.0112	-3.59440	0
4	0.32924	-4.4552	4.5718	-0.98880	0


## Exploratory Data Analysis

```
In [35]: 1 import seaborn as sns
        2 import matplotlib.pyplot as plt
        3 %matplotlib inline
```

**Create a Countplot of the Classes (Authentic 1 vs Fake 0)**

```
In [36]: 1 sns.countplot(x = 'Class', data = data)
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x1a25d57588>
```




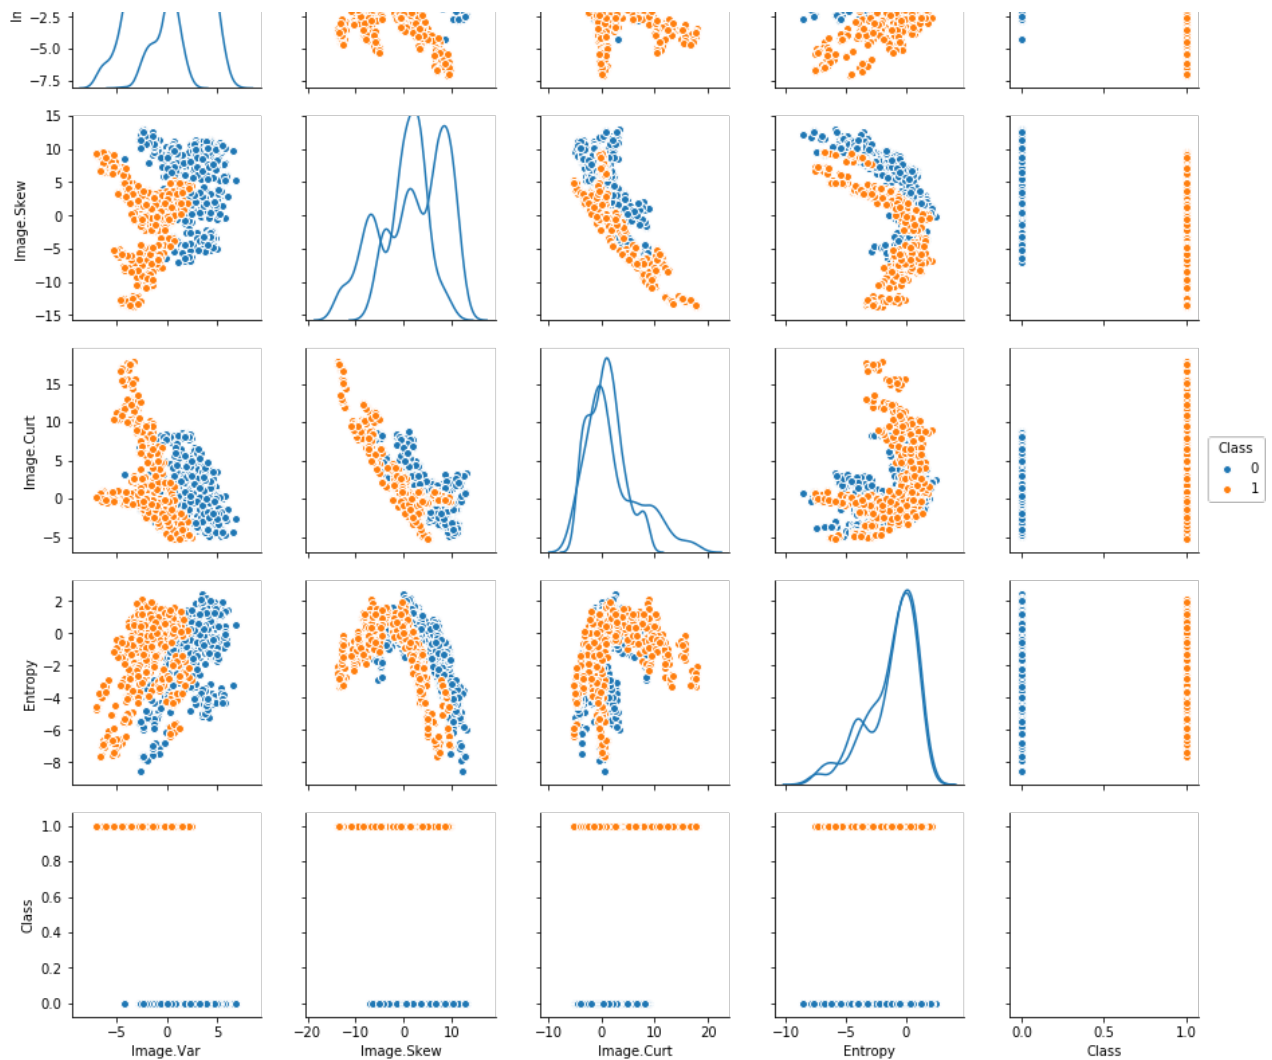
### Create a PairPlot of the Data with Seaborn, set Hue to Class

```
In [37]: 1 sns.pairplot(data = data, hue = "Class", diag_kind="kde")
```

```
/Users/Jayashri/anaconda/lib/python3.6/site-packages/scipy/stats/stats
.py:1713: FutureWarning: Using a non-tuple sequence for multidimension
al indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`
. In the future this will be interpreted as an array index, `arr[np.ar
ray(seq)]`, which will result either in an error or a different result
.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
/Users/Jayashri/anaconda/lib/python3.6/site-packages/statsmodels/nonpa
rametric/kde.py:494: RuntimeWarning: invalid value encountered in true
_divide
    binned = fast_linbin(X,a,b,gridsize)/(delta*nobs)
/Users/Jayashri/anaconda/lib/python3.6/site-packages/statsmodels/nonpa
rametric/kdetools.py:34: RuntimeWarning: invalid value encountered in
double_scalars
    FAC1 = 2*(np.pi*bw/RANGE)**2
/Users/Jayashri/anaconda/lib/python3.6/site-packages/numpy/core/fromnu
meric.py:83: RuntimeWarning: invalid value encountered in reduce
    return ufunc.reduce(obj, axis, dtype, out, **passkwargs)
```

```
Out[37]: <seaborn.axisgrid.PairGrid at 0x1a27909160>
```





## Data Preparation

When using Neural Network and Deep Learning based systems, it is a good practice to Standardize the data.

## Standard Scaling

```
In [38]: 1 from sklearn.preprocessing import StandardScaler
```

Create a `StandardScaler()` object called `scaler`.

```
In [39]: 1 scaler = StandardScaler()
```

Fit scaler to the features.

```
In [40]: 1 scaler.fit(X=data.drop("Class", axis = 1))
```

```
Out[40]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

**Use the .transform() method to transform the features to a scaled version.**

```
In [43]: 1 scaled_features = scaler.transform(X=data.drop("Class",
2                                                    axis = 1))
3 scaled_features
```

```
Out[43]: array([[ 1.12180565,  1.14945512, -0.97597007,  0.35456135],
 [ 1.44706568,  1.06445293, -0.89503626, -0.12876744],
 [ 1.20780971, -0.77735215,  0.12221838,  0.61807317],
 ...,
 [-1.47235682, -2.62164576,  3.75901744, -0.75488418],
 [-1.40669251, -1.75647104,  2.552043  , -0.04315848],
 [-1.04712236, -0.43982168,  0.29861555,  1.1364645 ]])
```

**Convert the scaled features to a dataframe and check the head of this dataframe to make sure the scaling worked.**

```
In [51]: 1 df_feat = pd.DataFrame(data = scaled_features,
2                               columns = data.columns[:-1])
3 df_feat.head()
```

```
Out[51]:
```

	Image.Var	Image.Skew	Image.Curt	Entropy
0	1.121806	1.149455	-0.975970	0.354561
1	1.447066	1.064453	-0.895036	-0.128767
2	1.207810	-0.777352	0.122218	0.618073
3	1.063742	1.295478	-1.255397	-1.144029
4	-0.036772	-1.087038	0.736730	0.096587

## Train Test Split

**Create two objects X and y which are the scaled feature values and labels respectively.**

```
In [56]: 1 x = df_feat
2 y = data["Class"]
```

**Use SciKit Learn to create training and testing sets of the data, using test\_size = 30%**

```
In [57]: 1 from sklearn.model_selection import train_test_split
```

```
In [58]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y,  
2                                                    test_size=0.3)
```

## Tensorflow

Create a list of feature column objects using `tf.feature.numeric_column()`

```
In [61]: 1 import tensorflow as tf  
2 X.columns
```

```
Out[61]: Index(['Image.Var', 'Image.Skew', 'Image.Curt', 'Entropy'], dtype='object')
```

```
In [70]: 1 df_feat.columns
```

```
Out[70]: Index(['Image.Var', 'Image.Skew', 'Image.Curt', 'Entropy'], dtype='object')
```

```
In [71]: 1 image_var = tf.feature_column.numeric_column("Image.Var")  
2 image_skew = tf.feature_column.numeric_column("Image.Skew")  
3 image_curt = tf.feature_column.numeric_column("Image.Curt")  
4 entropy = tf.feature_column.numeric_column("Entropy")
```

```
In [73]: 1 feat_cols=[image_var, image_skew, image_curt, entropy]
```

Create an object called classifier which is a `DNNClassifier` from `learn`. Set it to have 2 classes and a [10,20,10] hidden unit layer structure:

```
In [74]: 1 classifier = tf.estimator.DNNClassifier(hidden_units = [10,20,10],
2                                             n_classes= 2,
3                                             feature_columns = feat_cols)
```

```
INFO:tensorflow:Using default config.
WARNING:tensorflow:Using temporary folder as model directory: /var/folders/lw/sc_sdzmj67jcw7r58sdqrt2r0000gq/T/tmpkq1_q9ob
INFO:tensorflow:Using config: {'_model_dir': '/var/folders/lw/sc_sdzmj67jcw7r58sdqrt2r0000gq/T/tmpkq1_q9ob', '_tf_random_seed': None, '_save_summary_steps': 100, '_save_checkpoints_steps': None, '_save_checkpoints_secs': 600, '_session_config': None, '_keep_checkpoint_max': 5, '_keep_checkpoint_every_n_hours': 10000, '_log_step_count_steps': 100, '_train_distribute': None, '_device_fn': None, '_service': None, '_cluster_spec': <tensorflow.python.training.server_lib.ClusterSpec object at 0x1a3f7d7fd0>, '_task_type': 'worker', '_task_id': 0, '_global_id_in_cluster': 0, '_master': '', '_evaluation_master': '', '_is_chief': True, '_num_ps_replicas': 0, '_num_worker_replicas': 1}
```

```
1 ** Now create a tf.estimator.pandas_input_fn that takes in
  X_train, y_train, batch_size and set shuffle=True. Set batch_size
  parameter to 20. **
```

```
In [75]: 1 input_func = tf.estimator.inputs.pandas_input_fn(x = X_train,
2                                                         y = y_train,
3                                                         batch_size = 20,
4                                                         shuffle = True)
```

**Now train classifier to the input function. Use steps=500.**

```
In [76]: 1 classifier.train(input_fn=input_func, steps = 500)
```

```
INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Create CheckpointSaverHook.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Saving checkpoints for 0 into /var/folders/lw/sc_sdzmj67jcw7r58sdqrt2r0000gq/T/tmpkq1_q9ob/model.ckpt.
INFO:tensorflow:loss = 14.076602, step = 1
INFO:tensorflow:Saving checkpoints for 48 into /var/folders/lw/sc_sdzmj67jcw7r58sdqrt2r0000gq/T/tmpkq1_q9ob/model.ckpt.
INFO:tensorflow:Loss for final step: 0.38162005.
```

```
Out[76]: <tensorflow.python.estimator.canned.dnn.DNNClassifier at 0x1a3f7d71d0>
```

```
■ ■ ■ ■ ■ ■ ■ ■
```

## Model Evaluation

Create another `pandas_input_fn` that takes in the `X_test` data for `x`. This one won't need any `y_test` info since we will be using this for the network to create its own predictions. Set `shuffle=False` since we don't need to shuffle for predictions.

```
In [77]: 1 pred_fn = tf.estimator.inputs.pandas_input_fn(x=X_test,
2                                                    batch_size = len(X_test)
3                                                    shuffle = False)
```

Use the `predict` method from the classifier model to create predictions from `X_test`

```
In [78]: 1 predictions = list(classifier.predict(input_fn = pred_fn))

INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from /var/folders/lw/sc_sdzmj67jc
w7r58sdqrt2r0000gq/T/tmpkq1_q9ob/model.ckpt-48
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
```

```
In [79]: 1 #Let us look at predictions
2 predictions
```

```
Out[79]: [{'class_ids': array([1]),
  'classes': array([b'1'], dtype=object),
  'logistic': array([0.9935614], dtype=float32),
  'logits': array([5.038985], dtype=float32),
  'probabilities': array([0.0064386, 0.9935614], dtype=float32)},
 {'class_ids': array([1]),
  'classes': array([b'1'], dtype=object),
  'logistic': array([0.997971], dtype=float32),
  'logits': array([6.198182], dtype=float32),
  'probabilities': array([0.002029, 0.997971], dtype=float32)},
 {'class_ids': array([0]),
  'classes': array([b'0'], dtype=object),
  'logistic': array([0.00843123], dtype=float32),
  'logits': array([-4.7673454], dtype=float32),
  'probabilities': array([0.99156874, 0.00843123], dtype=float32)},
 {'class_ids': array([0]),
  'classes': array([b'0'], dtype=object),
  'logistic': array([0.00111933], dtype=float32),
  'logits': array([-6.793903], dtype=float32),
  'probabilities': array([0.99888067, 0.00111933], dtype=float32)}]
```

It is a dictionary with key 'class\_ids' and we want to get the first value i.e. `array[1]` or `array[0]`

```
In [82]: 1 final_preds = []
          2 for pred in predictions:
          3     final_preds.append(pred['class_ids'][0])
          4
```

**Now create a classification report and a Confusion Matrix.**

```
In [83]: 1 from sklearn.metrics import classification_report, confusion_matrix
```

```
In [84]: 1 print(confusion_matrix(y_test, final_preds))
```

```
[[224   7]
 [  0 181]]
```

```
In [85]: 1 print(classification_report(y_test, final_preds))
```

	precision	recall	f1-score	support
0	1.00	0.97	0.98	231
1	0.96	1.00	0.98	181
avg / total	0.98	0.98	0.98	412

## Comparison with Random Forest Classifier

The results from the DNN model seem to be accurate. Let's compare this to a Random Forest Classifier for a reality check!

Use SciKit Learn to Create a Random Forest Classifier and compare the confusion matrix and classification report to the DNN model

```
In [86]: 1 from sklearn.ensemble import RandomForestClassifier
```



```
In [94]: 1 rfc = RandomForestClassifier(n_estimators=200)
         2 rfc
```

```
Out[94]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

```
In [95]: 1 rfc.fit(X_train, y_train)
```

```
Out[95]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

```
In [96]: 1 rfc_pred = rfc.predict(X_test)
```

```
In [97]: 1 print(classification_report(y_test, rfc_pred))
```

	precision	recall	f1-score	support
0	1.00	0.99	0.99	231
1	0.99	0.99	0.99	181
avg / total	0.99	0.99	0.99	412

```
In [98]: 1 print(confusion_matrix(y_test, rfc_pred))
```

```
[[229  2]
 [ 1 180]]
```

**It looks like Random Forest Classifier does a very good job too and it is comparable to Deep Neural Networks.**

```
In [ ]: 1
```

