



(<http://www.pieriandata.com>)

Text Classification Assessment

This assessment is very much like the Text Classification Project we just completed, and the dataset is very similar.

The **moviereviews2.tsv** dataset contains the text of 6000 movie reviews. 3000 are positive, 3000 are negative, and the text has been preprocessed as a tab-delimited file. As before, labels are given as `pos` and `neg`.

We've included 20 reviews that contain either `NaN` data, or have strings made up of whitespace.

For more information on this dataset visit <http://ai.stanford.edu/~amaas/data/sentiment/>
(<http://ai.stanford.edu/~amaas/data/sentiment/>)

Task #1: Perform imports and load the dataset into a pandas DataFrame

For this exercise you can load the dataset from `'../TextFiles/moviereviews2.tsv'`.

```
In [1]: import pandas as pd
import numpy as np

df = pd.read_csv('../TextFiles/moviereviews2.tsv', sep='\t')
df.head()
```

Out[1]:

	label	review
0	pos	I loved this movie and will watch it again. Or...
1	pos	A warm, touching movie that has a fantasy-like...
2	pos	I was not expecting the powerful filmmaking ex...
3	neg	This so-called "documentary" tries to tell tha...
4	pos	This show has been my escape from reality for ...

```
In [2]: len(df)
```

Out[2]: 6000

Task #2: Check for missing values:

```
In [3]: # Check for NaN values:
df.isnull().sum()
```

Out[3]: label 0
review 20
dtype: int64

```
In [4]: df.dropna(inplace = True)
```

```
In [5]: df.isnull().sum()
```

Out[5]: label 0
review 0
dtype: int64

Task #3: Remove whitespace values:

```
In [11]: blanks = []

for i, lb, rv in df.itertuples():
    if rv.isspace():
        blanks.append(i)

print(len(blanks))
```

0

Task #4: Take a quick look at the `label` column:

```
In [12]: df['label'].head()
```

```
Out[12]: 0    pos
         1    pos
         2    pos
         3    neg
         4    pos
         Name: label, dtype: object
```

```
In [23]: df['label'].value_counts()
```

```
Out[23]: pos      2990
         neg      2990
         Name: label, dtype: int64
```

Task #5: Split the data into train & test sets:

You may use whatever settings you like. To compare your results to the solution notebook, use `test_size=0.33, random_state=42`

```
In [13]: from sklearn.model_selection import train_test_split
```

```
In [16]: X = df['review']
         y = df['label']

         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
```

Task #6: Build a pipeline to vectorize the data, then train and fit a model

You may use whatever model you like. To compare your results to the solution notebook, use `LinearSVC`.

```
In [18]: from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.svm import LinearSVC

text_clf = Pipeline([('tfidf', TfidfVectorizer()), ('clf', LinearSVC())]

text_clf.fit(X_train, y_train)
```

```
Out[18]: Pipeline(memory=None,
      steps=[('tfidf', TfidfVectorizer(analyzer='word', binary=False, d
encode_error='strict',
      dtype=<class 'numpy.float64'>, encoding='utf-8', input='conten
t',
      lowercase=True, max_df=1.0, max_features=None, min_df=1,
      ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=T
rue,...ax_iter=1000,
      multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
      verbose=0))])
```

Task #7: Run predictions and analyze the results

```
In [19]: # Form a prediction set
predictions = text_clf.predict(X_test)
```

```
In [20]: # Report the confusion matrix
from sklearn.metrics import confusion_matrix

print(confusion_matrix(y_test, predictions))
```

```
[[900  91]
 [ 63 920]]
```

```
In [21]: # Print a classification report
from sklearn.metrics import classification_report

print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
neg	0.93	0.91	0.92	991
pos	0.91	0.94	0.92	983
micro avg	0.92	0.92	0.92	1974
macro avg	0.92	0.92	0.92	1974
weighted avg	0.92	0.92	0.92	1974

```
In [22]: # Print the overall accuracy
from sklearn.metrics import accuracy_score

print(accuracy_score(y_test, predictions))
```

0.9219858156028369

Great job!